Building GPU-enabled Kubernetes single node cluster for MLOps experiments with new+second hand components

1. **SUPERMICRO MBD-X11SAE-M-O Micro ATX Server Motherboard LGA 1151 Intel C236** from newegg.com - $200 (Link to Newegg)

2. **Intel Xeon E3-1275 V6 Kaby Lake 3.8 GHz (4.2 GHz Turbo) LGA 1151 Server Processor Intel HD Graphics P630** - $360 (Link to Newegg)

3. 64GB of RAM: four DIMMs of **16GB DDR4-2400 UDIMM 1.2V CL17 (SKU: CT16G4DFD824A)** from Crucial: $250 (Link to Crucial)

4. **Sabrent Rocket Q 1TB NVMe PCIe M.2 2280 Internal SSD High Performance Solid State Drive R/W 3200/2000MB/s (SB-RKTQ-1TB)**: $110 (Link to Amazon)

5. **Supermicro SNK-P0046A4 Heatsink 2U+ Active Heatsink LGA1156 & LGA1155**: $25 (Link to Supermicro store)

6. **Nvidia Tesla M40 12GB** (Second hand): $125(Link to Ebay)

7. Dual 8 to 8 Graphics Power Cable (SKU: 030-0571-000): $6 (Link to Ebay)

Really old components:

8. HDD and SSD: Western Digital HDD 250GB SATA3 6.0Gb/s 7.2K (Used): $0, SSD 120 GB SATA3 6.0 Gb/s 2.5” (Used): $0, SSD 120 GB SATA3 6.0 Gb/s 2.5” (Used): $0

9. SuperMicro Chassis 733I-500B with 500W PSU (Used): $0 (Link to SuperMicro)

Final configuration: 1 x Xeon E3-1275, 64GB RAM, 1.5TB storage with 1TB on NVMe, NVIDIA GPU (Maxwell architecture) with 12GB GDDR5 for **total: $1076.0 (Ouch!)**

Additional Notes:

The 500W power supply coming with the mid tower chassis seems adequate for the graphics card provided you do not plan to reach frequently the max 250W which the card may need in peak memory and GPU utilization. The CPU drains no more than 73W, and the motherboard with the chipset and 64GB of RAM. Note that the chassis has an extension space allowing mounting larger power supply in case the one with 500W is not sufficient.

Use Ubuntu 20.04 and upgrade to 20.10.

Connect all disks (NVMe, SSD1, SSD2, and HDD) in a single file system using LVM2

Do not configure swap partition in case you are going to run Kubernetes.